


CS 498 HW 6 Yu Che Wang/ yuchecw2

1. All the points (row number) I removed (indexed on the original dataset) as outlier points by looking at the influence plot:


364, 365, 368, 370, 371, 372, 380, 405, 410, 418.

I remove those points with high leverage, standardized residuals, and cook distance.


The radius of the data point represents its cook distance.




2. By using `scipy.stats.boxcox`, we obtain the best lambda value = 0.219.



3. Diagnostic plot before removing outliers



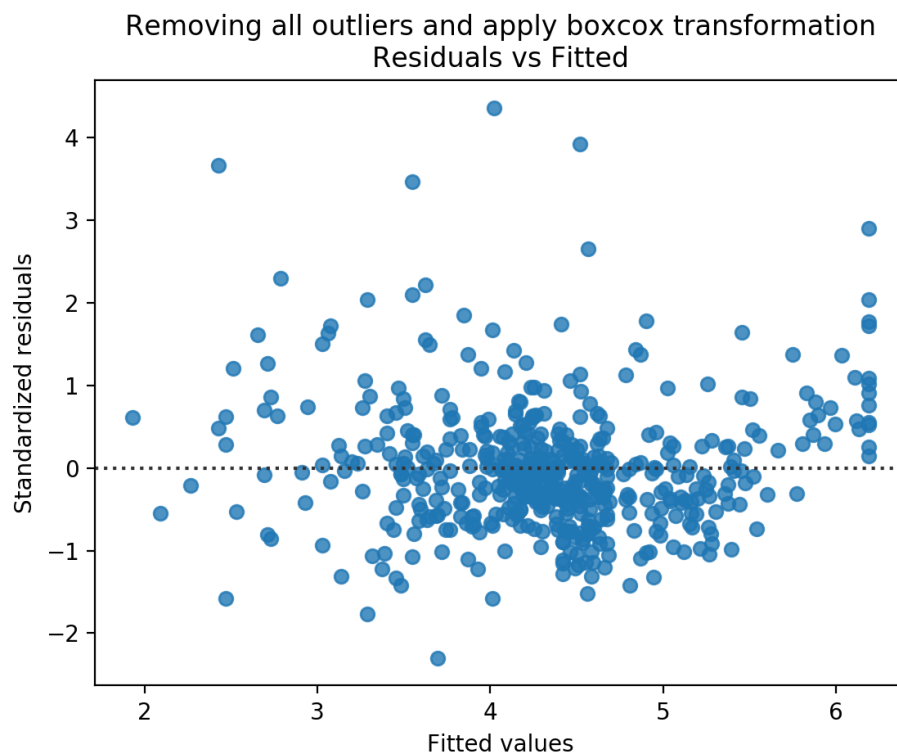
Diagnostic plot after removing outliers



4. Raw data

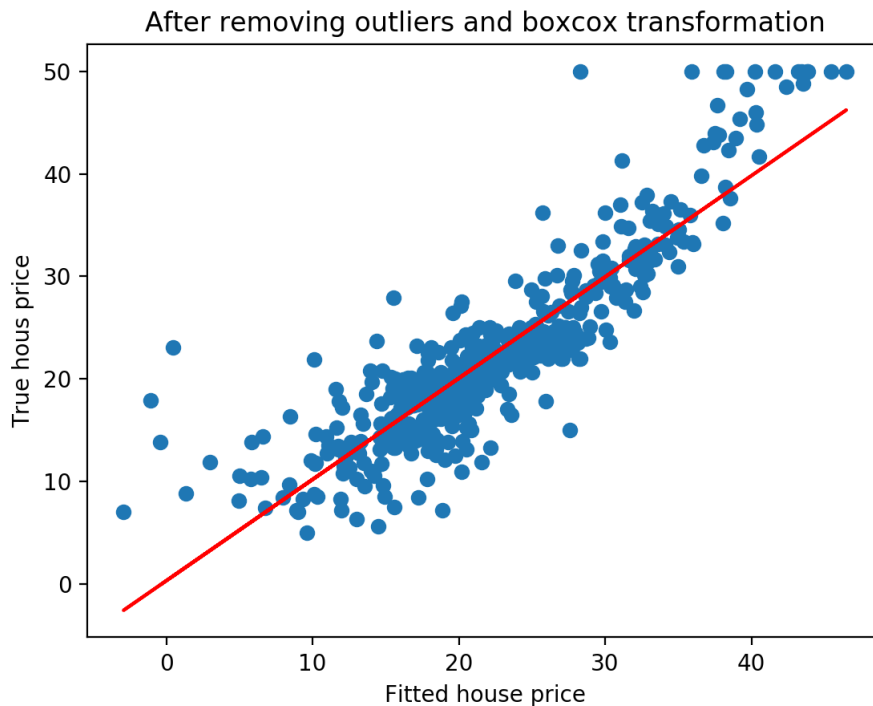


5. Removing all outliers and transforming the dependent variable



6. In the first plot, there are some outliers, data that are 4-6 standardized residuals away from the mean. Most of the data points in the second plot are less than two standardized residuals away from the mean. In addition, the data points are evenly distributed across the dotted line.

7. Fitted house price vs True house price



The data points lie on a line, which means that after removing outliers and applying box-cox transformation, the residuals are nearly zero.

8. Code fragments

Linear regression

```
15 # Linear regression
16 X = data[:,13]
17 y = data[:,13].reshape(-1,1)
18 # Method 1
19 lm = sm.OLS(y, sm.add_constant(X)).fit()
20 # Method 2
21 reg = LinearRegression().fit(X, y)
22 r_squared = reg.score(X,y)
```

Box-cox transformation

```
132 # Boxcox transformation
133 y_boxcox, maxlog = stats.boxcox(y_remove)
```

Box-cox transformation with different values

```
96 n_interval = 50
97 lambda_list = np.linspace(-2, 2, n_interval)
98 llf_list = np.zeros(n_interval)
99 for i in range(lambda_list.shape[0]):
100     llf_list[i] = stats.boxcox_llf(lambda_list[i], y_remove)
```

To apply box-cox transformation with a specific value λ , one can use

$y_{\text{boxcox}} = \text{stats.boxcox}(y, \text{lambda}=\lambda)$

9. Entire code

```
1 import numpy as np
2 import numpy.linalg as la
3 import pandas as pd
4 import seaborn as sns
5 import scipy.stats as stats
6 import matplotlib.pyplot as plt
7 from sklearn.linear_model import LinearRegression
8 from statsmodels.stats.outliers_influence import OLSInfluence
9 import statsmodels.api as sm
10 import statsmodels.formula.api as smf
11 # Load the data
12 data = np.loadtxt('data.txt')
13 df = pd.DataFrame(data)
14
15 # Linear regression
16 X = data[:, :13]
17 y = data[:, 13].reshape(-1, 1)
18 # Method 1
19 lm = sm.OLS(y, sm.add_constant(X)).fit()
20 # Method 2
21 reg = LinearRegression().fit(X, y)
22 r_squared = reg.score(X, y)
23
24 # Generate influence plot
25 fig, ax = plt.subplots(figsize=(9, 6))
26 fig = sm.graphics.influence_plot(lm, alpha=0.001, ax=ax, criterion='cooks')
27 fig.show()
28
29 #####
30 # Calculate leverage
31 #####
32 beta_hat = np.dot(la.inv(np.dot(X.T, X)), np.dot(X.T, y))
33 hat_matrix = X@la.inv(np.dot(X.T, X))@(X.T)
34 # plt.bar(list(range(hat_matrix.shape[0])), np.diag(hat_matrix))
35 # plt.boxplot(np.diag(hat_matrix))
36 leverage = np.diag(hat_matrix)
37 # Remove leverage outliers and record their indices
38 outliers_num_leverage = []
39 for i in range(leverage.shape[0]):
40     if (leverage[i] > np.percentile(leverage, 75) + 1.5 * stats.iqr(leverage)):
41         outliers_num_leverage.append(i)
42 leverage_remove_outliers = leverage[leverage <= np.percentile(leverage, 75) + 1.5 * stats.iqr(leverage)]
43
44 # Calculate residuals and mean square error
45 y_predicted = X@beta_hat # Fitted value
46 e = y - X@beta_hat
47 N = y.shape[0]
48 mean_squared_error = np.dot(e.T, e)/N
49 # r_squared = np.var(X@beta_hat)/np.var(y)
50
51 #####
52 # Calculate the cook distance for each data
53 #####
54 cook_distance = np.zeros((y.shape[0], 1))
55 for i in range(y.shape[0]):
56     X_remove = np.delete(X, i, 0)
57     y_remove = np.delete(y, i, 0)
58     beta_i = np.dot(la.inv(np.dot(X_remove.T, X_remove)), np.dot(X_remove.T, y_remove))
59     y_p = X_remove@la.inv(np.dot(X_remove.T, X_remove))@(X_remove.T@y_remove)
60     cook_distance[i] = np.dot((y_p - X_remove@beta_i).T, y_p - X_remove@beta_i)/(N * mean_squared_error)
61 # plt.bar(list(range(y.shape[0])), cook_distance)
62 # plt.boxplot(cook_distance)
63 outliers_num_cook_distance = []
64 for i in range(cook_distance.shape[0]):
65     if (cook_distance[i] > np.percentile(cook_distance, 75) + 1.5 * stats.iqr(cook_distance)):
66         outliers_num_cook_distance.append(i)
67 cook_distance_remove_outliers = cook_distance[cook_distance <= np.percentile(cook_distance, 75) + 1.5 * stats.iqr(cook_distance)]
68
69 #####
70 # Calculate standardized residuals
71 #####
72 s = np.zeros((y.shape[0], 1))
73 for i in range(y.shape[0]):
74     s[i] = e[i]/np.sqrt(mean_squared_error*(1-leverage[i]))
75 # plt.bar(list(range(y.shape[0])), np.abs(s))
76 # plt.boxplot(s)
77 outliers_num_s = []
78 for i in range(s.shape[0]):
79     if (s[i] > np.percentile(s, 75) + 1.5 * stats.iqr(s)):
80         outliers_num_s.append(i)
81
```

```

82 # Remove outliers
83 by_eye = True
84 if by_eye:
85     outliers_num = [364, 365, 368, 370, 371, 372, 380, 405, 410, 418]
86 else:
87     outliers_num = np.union1d(outliers_num_leverage, outliers_num_cook_distance)
88     outliers_num = np.union1d(outliers_num, outliers_num_s)
89
90 lm_remove = sm.OLS(y_remove, sm.add_constant(X_remove)).fit()
91 X_remove = np.delete(X, outliers_num, 0)
92 y_remove = np.delete(y, outliers_num, 0)
93 s_remove = np.delete(s, outliers_num, 0)
94
95 # Page 1
96 n_interval = 50
97 lambda_list = np.linspace(-2, 2, n_interval)
98 llf_list = np.zeros(n_interval)
99 for i in range(lambda_list.shape[0]):
100     llf_list[i] = stats.boxcox_llf(lambda_list[i], y_remove)
101 fig, ax = plt.subplots(figsize=(8,4))
102 plt.plot(lambda_list, llf_list)
103 plt.xlabel('Lambda value')
104 plt.ylabel('Log likelihood')
105 plt.title('Box-Cox Transformation Curve')
106 plt.show()
107
108 # Page 2
109 # Generate diagnostic plots before removing outliers
110 before = True # Change this variable to False to generate diagnostic plot after removing outliers
111 if before:
112     model_residuals = lm.resid
113     model_norm_residuals = lm.get_influence().resid_studentized_internal
114     model_leverage = lm.get_influence().hat_matrix_diag
115     model_cooks = lm.get_influence().cooks_distance[0]
116
117     diagnostic_plot = plt.figure()
118     plt.scatter(model_leverage, model_norm_residuals, alpha=0.5)
119     sns.regplot(model_leverage, model_norm_residuals, scatter=False, ci=False, lowess=True,
120                 line_kws={'color':'red', 'lw':1, 'alpha':0.8}, label='Cook's distance')
121     diagnostic_plot.axes[0].set_xlim(0, max(model_leverage)+0.01)
122     diagnostic_plot.axes[0].set_ylim(-3,5)
123     plt.title('Residuals vs Leverage\nBefore removing outliers')
124     plt.xlabel('Leverage')
125     plt.ylabel('Standardized Residuals')
126     plt.legend()
127
128     leverage_top_3 = np.flip(np.argsort(model_cooks), 0)[:3]
129     for i in leverage_top_3:
130         diagnostic_plot.axes[0].annotate(i, xy=(model_leverage[i], model_norm_residuals[i]))
131     plt.show()
132
133 # Boxcox transformation
134 y_boxcox, maxlog = stats.boxcox(y_remove)
135
136 # Influence plots after removing outliers and boxcox transformation
137 lm_remove = sm.OLS(y_boxcox, sm.add_constant(X_remove)).fit()
138 fig, ax = plt.subplots(figsize=(9,6))
139 fig = sm.graphics.influence_plot(lm_remove, alpha=0.001, ax=ax, criterion='cooks')
140 fig.show()
141
142 # Page 3
143 # Standardized residuals vs Fitted values without any transformation
144 sns.residplot(y_predicted, s)
145 plt.xlabel('Fitted values')
146 plt.ylabel('Standardized residuals')
147 plt.title('Without any transformation\nResiduals vs Fitted')
148 plt.show()
149 # Standardized residuals vs Fitted values removing all outliers and boxcox transformation
150 beta_hat_prime = np.dot(la.inv(np.dot(X_remove.T, X_remove)), np.dot(X_remove.T, y_remove))
151 y_predicted_prime = X_remove@beta_hat_prime # Fitted house price
152 sns.residplot(y_boxcox, np.delete(s, outliers_num))
153 plt.xlabel('Fitted values')
154 plt.ylabel('Standardized residuals')
155 plt.title('Removing all outliers and apply boxcox transformation\nResiduals vs Fitted')
156 plt.show()
157
158 # Page 4
159 # Fitted house price vs True house price
160 plt.scatter(y_predicted_prime, y_remove)
161 plt.plot(y_predicted_prime, LinearRegression().fit(y_predicted_prime, y_remove).predict(y_predicted_prime), 'r')
162 plt.xlabel('Fitted house price')
163 plt.ylabel('True house price')
164 plt.title('After removing outliers and boxcox transformation')
165 plt.show()
166

```