

HW11: Variational Autoencoders

Yu Che Wang / yuchecw2

December 11, 2018

Abstract

Apply variational autoencoders to MNIST dataset using TensorFlow framework.

1 Single MNIST digits

For 10 pairs of MNIST test images of the same digit, selected at random, compute the code for each image of the pair. Now compute 7 evenly spaced linear interpolates between these codes, and decode the result into images.




Figure 1: Single MNIST digits.

2 Different MNIST digits

For 10 pairs of MNIST test images of different digits, selected at random, compute the code for each image of the pair. Now compute 7 evenly spaced linear interpolates between these codes, and decode the result into images.




Figure 2: Different MNIST digits.

3 Codes

```
1 ''' Autoencoder for MNIST dataset'''
2 from __future__ import division, print_function, absolute_import
3
4 import tensorflow as tf
5 from tensorflow.examples.tutorials.mnist import input_data
6 import numpy as np
7 import matplotlib.pyplot as plt
8 import argparse
9
10 def weight_variable(shape, name):
11     return tf.Variable(tf.truncated_normal(shape, stddev=0.1), name=name)
12
13 def bias_variable(shape, name):
14     return tf.Variable(tf.constant(0.1, shape=shape), name=name)
15
16 def encoder(x, in_channels, hid_channels_1, hid_channels_2, name='encoder'):
17     with tf.name_scope(name):
18         w1 = weight_variable([in_channels, hid_channels_1], name='w1')
19         b1 = bias_variable([hid_channels_1], name='b1')
20         hid_out_1 = tf.nn.sigmoid(tf.matmul(x, w1) + b1)
21         w2 = weight_variable([hid_channels_1, hid_channels_2], name='w2')
22         b2 = bias_variable([hid_channels_2], name='b2')
23         hid_out_2 = tf.nn.sigmoid(tf.matmul(hid_out_1, w2) + b2)
24         return hid_out_2
25
26 def decoder(x, hid_channels_2, hid_channels_1, out_channels, name='decoder'):
27     with tf.name_scope(name):
28         w1 = weight_variable([hid_channels_2, hid_channels_1], name='w1')
29         b1 = bias_variable([hid_channels_1], name='b1')
30         hid_out_1 = tf.nn.sigmoid(tf.matmul(x, w1) + b1)
31         w2 = weight_variable([hid_channels_1, out_channels], name='w2')
32         b2 = bias_variable([out_channels], name='b2')
33         hid_out_2 = tf.nn.sigmoid(tf.matmul(hid_out_1, w2) + b2)
34         return hid_out_2
```

Figure 3: Helper functions.

```
36 def train():
37     mnist = input_data.read_data_sets('/tmp/data', one_hot=True)
38     img_test = mnist.test.images
39     label_test = mnist.test.labels
40     label_img_dict = {}
41     for i in range(label_test.shape[0]):
42         num = np.argmax(label_test[i,:])
43         if num not in label_img_dict.keys():
44             label_img_dict[num] = img_test[i,:]
45         else:
46             label_img_dict[num] = np.vstack((label_img_dict[num], img_test[i,:]))
47
48     # Network parameters
49     in_channels = 784 # = out_channels
50     hid_channels_1 = 256
51     hid_channels_2 = 32
52     # Training parameters
53     learning_rate = 0.01
54     num_epochs = 100000
55     display_epoch = 1000
56     batch_size = 1
57
58     x = tf.placeholder(tf.float32, shape=[None, in_channels], name='x') # Input image a.k.a ground truth
59     y = tf.placeholder(tf.float32, shape=[None, hid_channels_2], name='y') # Encoder output
60
61     # Auto-encoder Model
62     encoder_out = encoder(x, in_channels, hid_channels_1, hid_channels_2) # Latent representation
63     decoder_out = decoder(y, hid_channels_2, hid_channels_1, in_channels) # Reconstructed image
64
65     # Calculate loss
66     loss = tf.losses.mean_squared_error(x, decoder_out)
67     optimizer = tf.train.RMSPropOptimizer(learning_rate).minimize(loss)
68
69     # Initialize the variables
70     init = tf.global_variables_initializer()
```

Figure 4: Main function.

```

72 with tf.Session() as sess:
73     sess.run(init)
74     print('Start session!')
75     # Training
76     loss_list = []
77     for epoch in range(num_epochs+1):
78         batch = mnist.train.next_batch(batch_size)
79         en_out = sess.run(encoder_out, feed_dict={x: batch[0]})
80         _, mse_loss = sess.run([optimizer, loss], feed_dict={x: batch[0], y: en_out})
81         loss_list.append(mse_loss)
82         if epoch % display_epoch == 0:
83             print('Epoch %d, loss: %.3f'%(epoch, mse_loss))
84     plt.figure()
85     plt.plot(list(range(len(loss_list))), loss_list)
86     plt.xlabel('Epoch number')
87     plt.ylabel('Loss')
88     plt.title('Loss v.s. Epoch')
89     plt.show()
90     # Testing
91     for are_same_digits in [True, False]:
92         plt.figure(figsize=(20,10))
93         for digit in range(10):
94             digit1, digit2 = None, None
95             if are_same_digits: # Same digits
96                 digit_arr = label_img_dict[digit]
97                 # Randomly select two same digits
98                 same_digits = digit_arr[np.random.choice(digit_arr.shape[0], 2, replace=False), :]
99                 digit1 = same_digits[0,:].reshape(1,-1)
100                digit2 = same_digits[1,:].reshape(1,-1)
101            else: # Different digits
102                # Randomly select two different digits
103                permutation_arr = np.random.permutation(10)
104                num1 = permutation_arr[0]
105                num2 = permutation_arr[1]
106                digit1 = label_img_dict[num1][np.random.choice(label_img_dict[num1].shape[0], 1, replace=False), :]
107                digit2 = label_img_dict[num2][np.random.choice(label_img_dict[num2].shape[0], 1, replace=False), :]
108            # Calculate corresponding codes
109            code1 = sess.run(encoder_out, feed_dict={x: digit1})
110            code2 = sess.run(encoder_out, feed_dict={x: digit2})
111            reconstructed_img = [None]*9
112            reconstructed_img[0], reconstructed_img[8] = digit1, digit2
113            # Compute 7 evenly spaced linear interpolates
114            for t in range(1, 8):
115                code = ((8-t)/8)*code1 + (t/8)*code2
116                reconstructed_img[t] = sess.run(decoder_out, feed_dict={y: code})
117            for idx in range(9):
118                plt.subplot(10,9,9*digit+idx+1)
119                plt.imshow(reconstructed_img[idx].reshape(28,28), cmap='gray')
120        if are_same_digits:
121            plt.savefig('same.png')
122        else:
123            plt.savefig('different.png')

```

Figure 5: (Continued)Main function.